

Case Study #4: Long-Term Return-on-Investment at Microsoft via Short-Term Proxy

ALICE Team, MSR New England

Maggie Hei



Microsoft



EconML

Session Goals

- Understand the common challenges of learning **holistic ROI** from both business and technical perspective
- Introduce a **unified pipeline** to estimate the long-term effect of multiple investments from observational data in a high dimensional manner
- Learn how to use **EconML** Python package to solve this problem in a few lines of code



Outline

- Business Background
- Technical Challenges
- Methodology
- EconML Solution
- Wrap up



Business Background

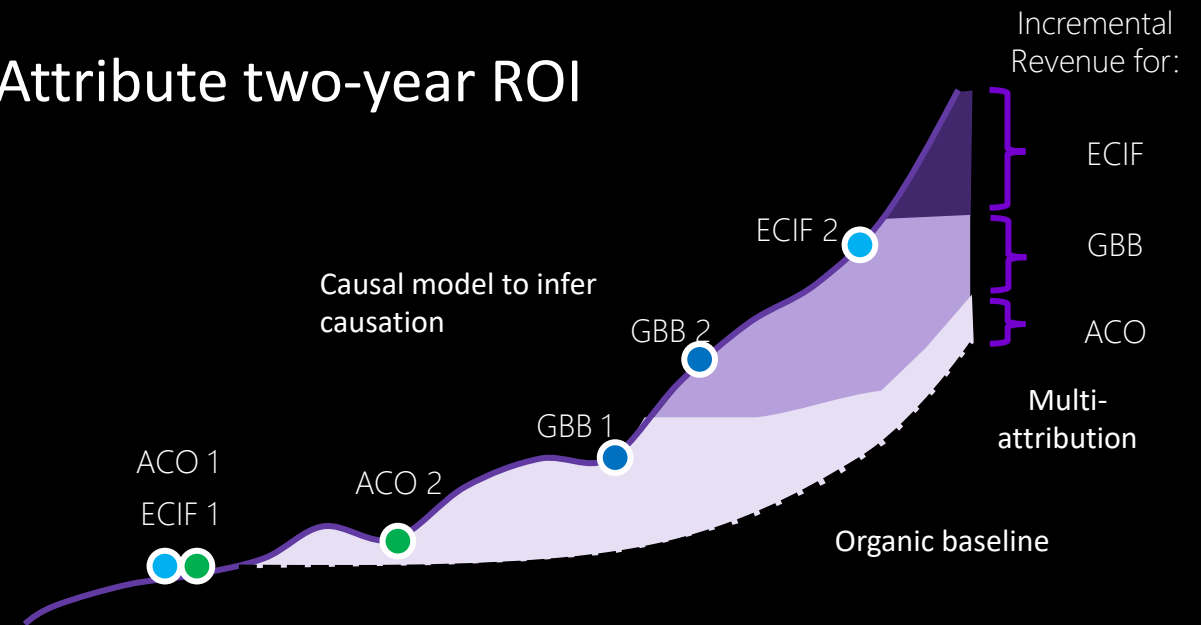
- Microsoft provides multiple monetary fundings to enterprise customers in support of products adoption. Which of these programs ("investment") are more successful than others?
- **Issues:**
 - Different models used to calculate Return on Investment separately for each program.
 - We want to estimate long-term success (e.g. two-year effect), but we can't wait that long to evaluate a program.
- **Goal:** Can we **attribute the long-term ROI** of multiple programs in a holistic manner with only **short-term data**?



Why Holistic ROI?

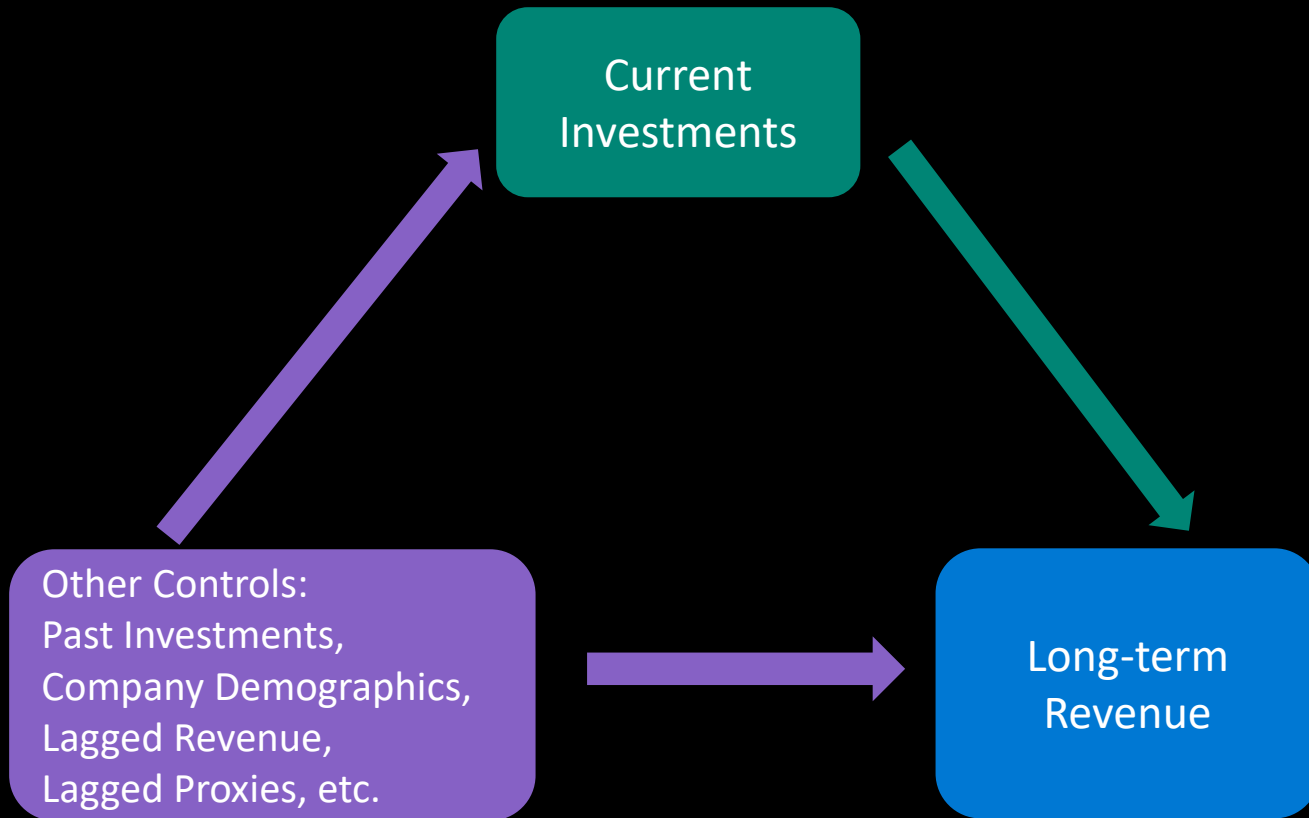
- Holistic ROI model gives us a way to assess the **relative impact of multiple investments across accounts, over time.**
 - Holistic view
 - Organic growth
 - Causal model
 - Attribution across investments

Attribute two-year ROI



Technical Challenge: Confounding Effect

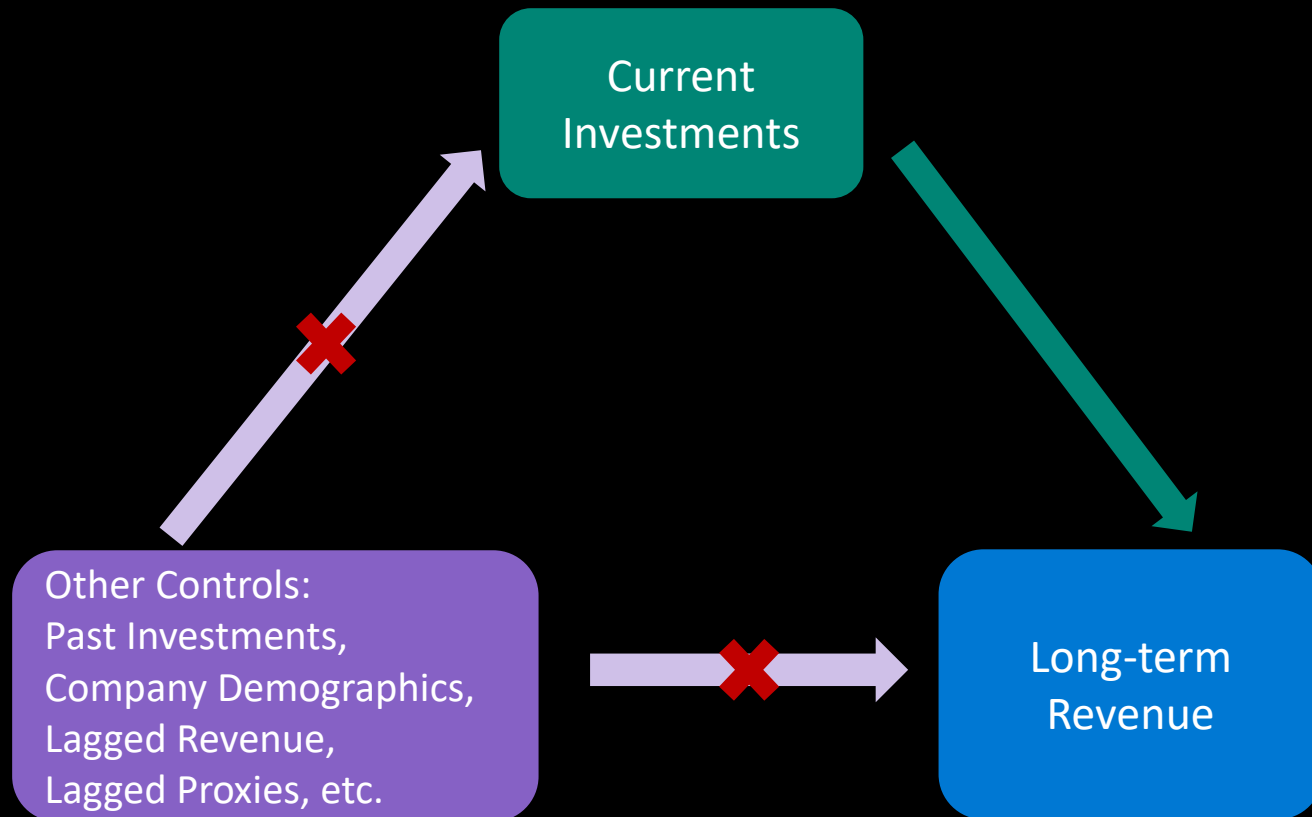
- Start with simple causal graph



- Hidden Factors** that affect both the target investment and the long-term revenue:
 - Larger customers might receive more investments
 - Customers would purchase some products even with no incentives
 - The effect might come from the past incentives

Methodology

- Careful causal modeling



- Double Machine Learning (DML)** nets out effects from confounders by fitting a two stage ML model
 - Allow high dimensional dataset
 - Construct confidence intervals

Only **direct effect** of target investments on future revenue remains!

Double Machine Learning

(Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins.
"Double/debiased machine learning for treatment and structural parameters")

$$Y = \theta(X) \times T + g(X, W) + \varepsilon$$
$$T = f(X, W) + \eta$$

1. Regress $Y \sim W$, learn \hat{Y}
 - **Predict future revenue** using past investments and current customer attributes
2. Regress $T \sim W$, learn \hat{T}
 - **Predict current investments** using past investments and current customer attributes
3. Linear regression on residuals: $(Y - \hat{Y}) \sim (T - \hat{T}) \otimes \phi(X)$
 - **Calculate residual**, "surprise" components of current investment and future revenue
4. $\theta(X) = \langle \vec{a}, \phi(x) \rangle$ where \vec{a} is the coefficient vector from the final regression
Note: for θ constant, θ is the coefficient from $(Y - \hat{Y}) \sim (T - \hat{T})$
 - **Correlation** of these residuals identifies average causal effect of investments

EconML Solution

```
# print data shape
print("Outcome shape: ", panelY.shape)
print("Treatment shape: ", panelT.shape)
print("Controls shape: ", panelX.shape)
```

```
Outcome shape: (5000, 4)
Treatment shape: (5000, 4, 3)
Controls shape: (5000, 4, 89)
```

```
# imports
from econml.dml import LinearDML
from sklearn.linear_model import LassoCV, MultiTaskLassoCV

# initiate DML estimator
est = LinearDML(
    model_y=LassoCV(max_iter=1000), # nuisance model y
    model_t=MultiTaskLassoCV(max_iter=1000), # nuisance model t
)
# fit treatment_0 on total revenue
est.fit(np.sum(panelY, axis=1), panelT[:, 0], X=None, W=panelX[:, 0])
# print treatment effect summary
est.summary(alpha=0.05)
```

Coefficient Results: X is None, please call intercept_inference to learn the constant!

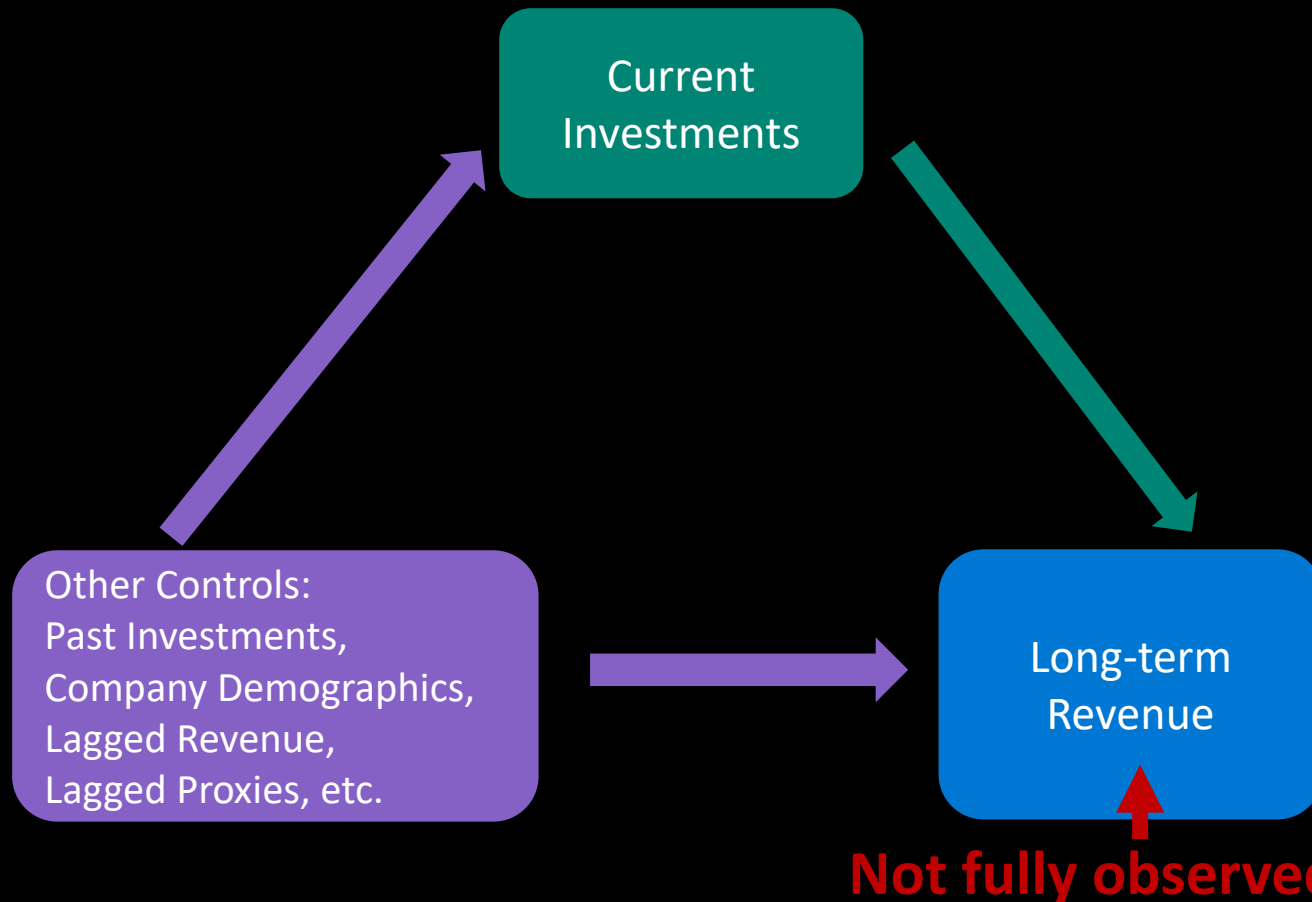
CATE Intercept Results

	point_estimate	stderr	zstat	pvalue	ci_lower	ci_upper
cate_intercept T0	2.927	0.206	14.219	0.0	2.523	3.33
cate_intercept T1	1.13	0.038	29.951	0.0	1.056	1.204
cate_intercept T2	0.528	0.058	9.135	0.0	0.415	0.642

In this case, $X = \text{None}$.
cate_intercept will be the
constant effect θ we want
to estimate!

Technical Challenge: Unobserved Outcome

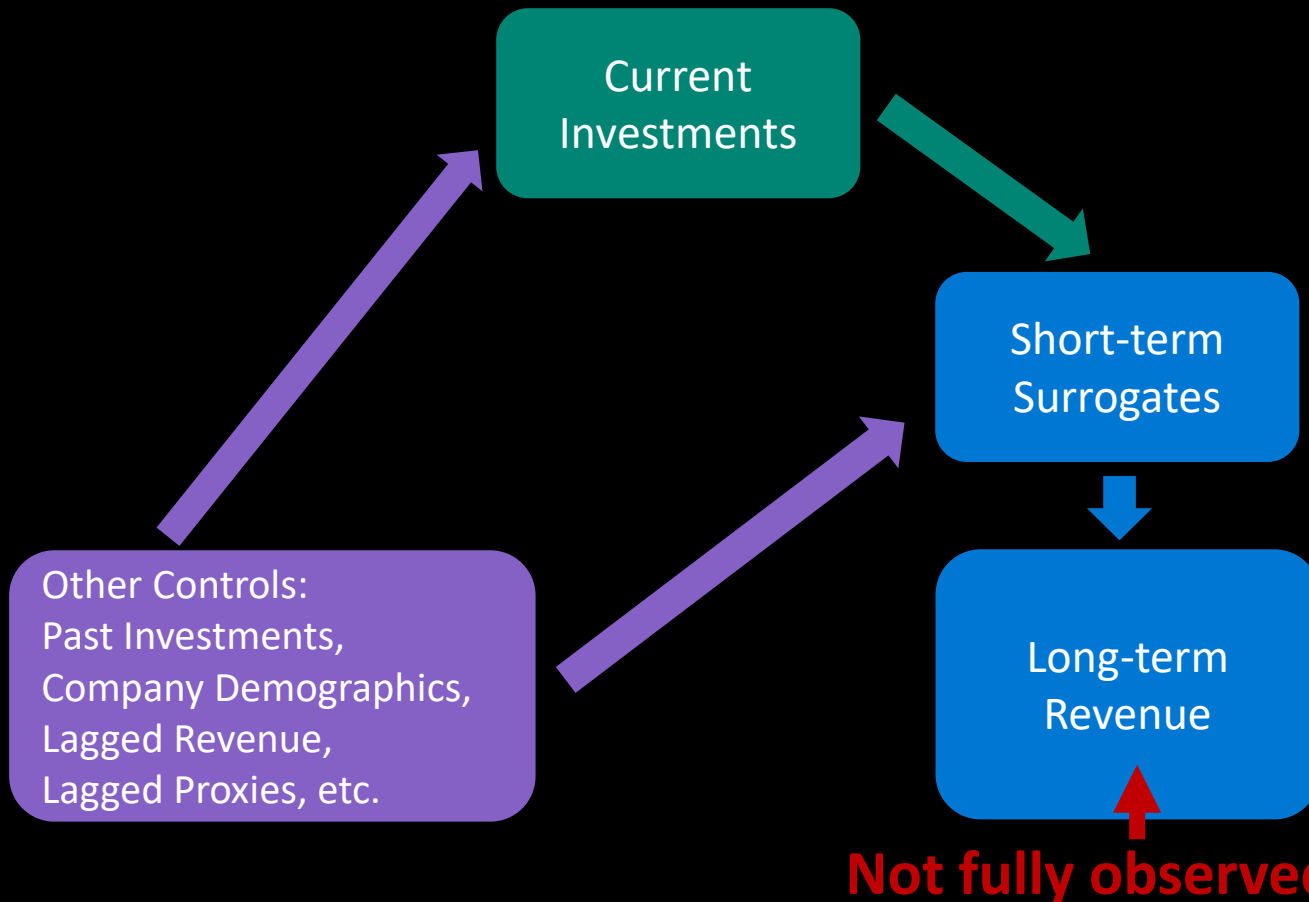
- Long-term revenue is not fully observed yet.



- Can we find **short-term observed surrogates** that are indicative of a customer's long-term revenue?

Technical Challenge: Unobserved Outcome

- Long-term revenue is not fully observed yet.




- Can we find **short-term observed surrogates** that are indicative of a customer's long-term revenue?

Methodology

- Estimation of Long-Term Effects with Surrogates

(Prentice, 1989; Begg & Leung, 2000; Frangakis & Rubin, 2002; Freedman et al., 1992; Athey et al., 2020)

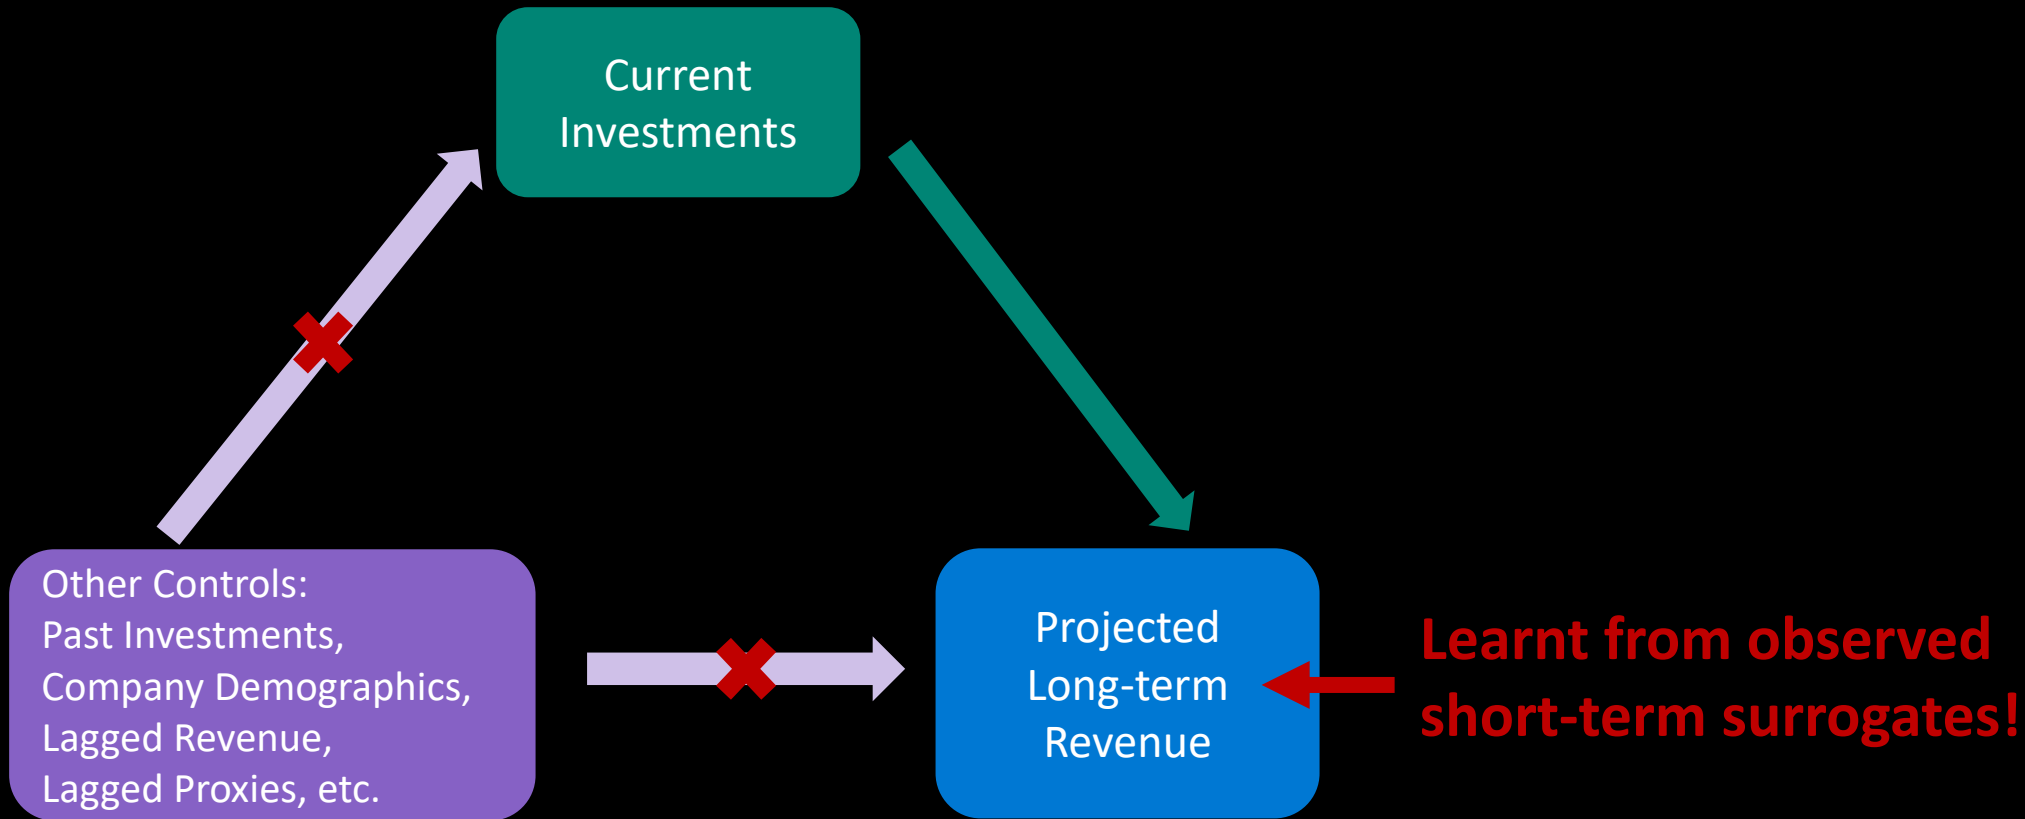


Assumption: The treatment effect of investments T on revenue Y could only go through surrogate S .

- Suppose we have a **historical dataset** (O) contains both long-term revenue Y and surrogates S , and a **current dataset** (E) contains only surrogates S .
- We can estimate $g(S) := E[Y|S]$ from O by running a regression $Y \sim S$ using any ML models.
- Then we could predict the current long-term $\hat{Y} = g(S)$ from E .
- \hat{Y} will be the projected long-term revenue, used as the outcome for DML model.

Updated Casual Graph

- Surrogates Index + DML



EconML Solution

Historical Data

Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2018	...	\$1,000	\$10,000
2	A	2019	...	\$2,000	\$12,000
3	A	2020	...	\$3,000	\$15,000
4	B	2018	...	\$0	\$5,000
5	B	2019	...	\$100	\$10,000
6	B	2020	...	\$1,200	\$7,000
7	C	2018	...	\$1,000	\$20,000
8	C	2019	...	\$1,500	\$25,000
9	C	2020	...	\$500	\$15,000

```
# train surrogate index
XS_0 = np.hstack([panelX_O[:, 0], panelY_O[:, :1]]) # concatenate controls and surrogates from historical dataset
XS_E = np.hstack([panelX_E[:, 0], panelY_E[:, :1]]) # concatenate controls and surrogates from current dataset
TotalY_0 = np.sum(panelY_O, axis=1) # total revenue from historical dataset
unadjusted_proxy_model = LassoCV().fit(XS_0, TotalY_0) # train proxy model from historical dataset
sindex = unadjusted_proxy_model.predict(XS_E) # predict current long term revenue
```

Current Data with New Investment

Company	Year	Features	Controls/Surrogates	T1	T2	T3	SurrRev
1	A	2021	...	\$1,000	\$10,000
2	B	2021	...	\$0	\$5,000
3	C	2021	...	\$2,000	\$15,000

EconML Solution

Historical Data

	Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2018	\$1,000	\$10,000
2	A	2019	\$2,000	\$12,000
3	A	2020	\$3,000	\$15,000
4	B	2018	\$0	\$5,000
5	B	2019	\$100	\$10,000
6	B	2020	\$1,200	\$7,000
7	C	2018	\$1,000	\$20,000
8	C	2020	\$1,500	\$25,000
9	C	2020	\$500	\$15,000

```
# train surrogate index
XS_O = np.hstack([panelX_O[:, 0], panelY_O[:, :1]]) # concatenate controls and surrogates from historical dataset
XS_E = np.hstack([panelX_E[:, 0], panelY_E[:, :1]]) # concatenate controls and surrogates from current dataset
TotalY_O = np.sum(panelY_O, axis=1) # total revenue from historical dataset
unadjusted_proxy_model = LassoCV().fit(XS_O, TotalY_O) # train proxy model from historical dataset
sindex = unadjusted_proxy_model.predict(XS_E) # predict current long term revenue
```

```
# Learn treatment effect on surrogate index
from econml.dml import LinearDML

est = LinearDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
# fit treatment_0 on total revenue from current dataset
est.fit(sindex, panelT_E[:, 0], X=None, W=panelX_E[:, 0])
# print treatment effect summary
est.summary(alpha=0.05)
```

Coefficient Results: X is None, please call intercept_inference to learn the constant!

CATE Intercept Results

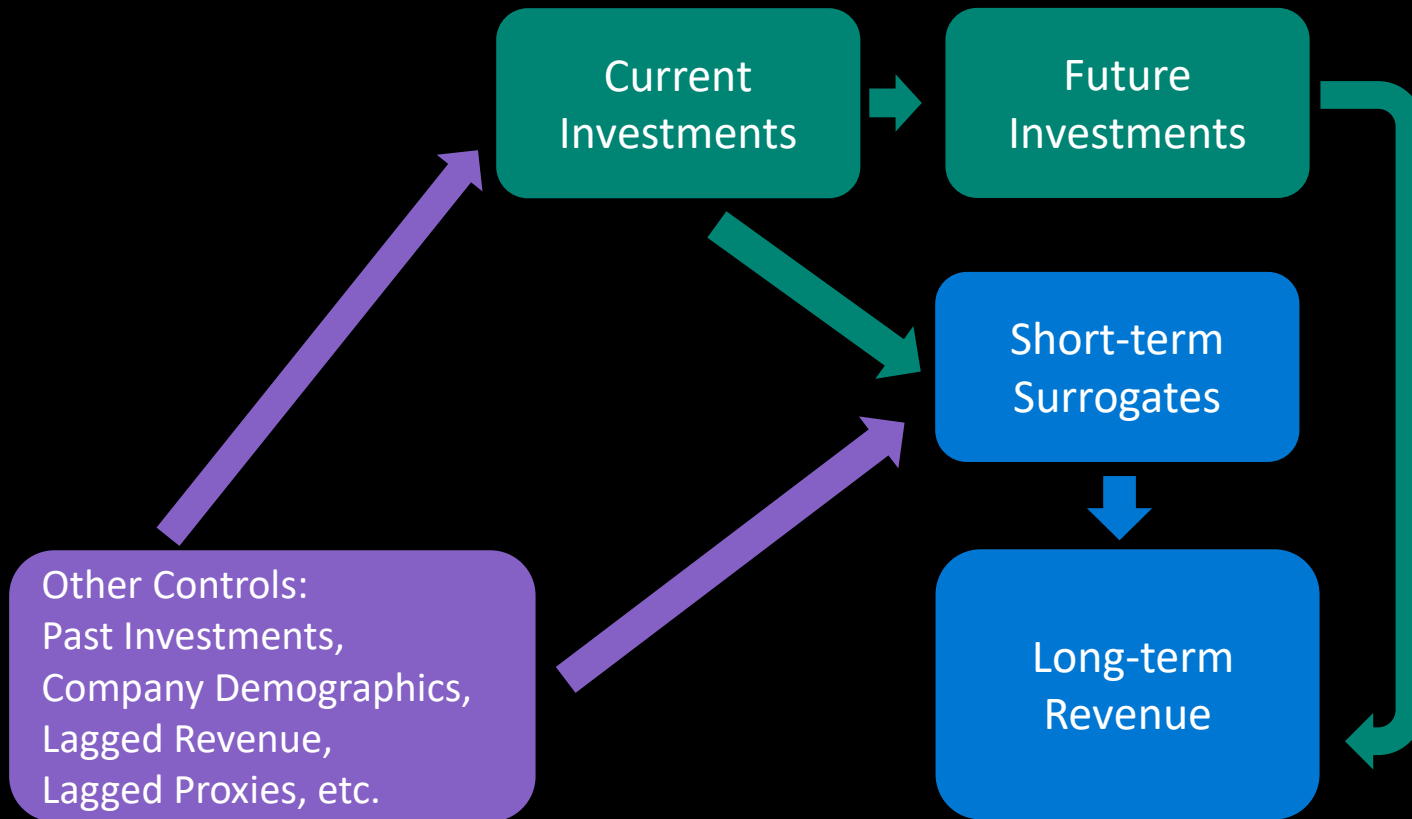
	point_estimate	stderr	zstat	pvalue	ci_lower	ci_upper
cate_intercept T0	1.415	0.03	46.428	0.0	1.356	1.475
cate_intercept T1	0.662	0.018	37.064	0.0	0.627	0.697
cate_intercept T2	0.261	0.006	44.417	0.0	0.249	0.272

Current Data with New Investment

	Company	Year	Features	Controls/Surrogates	T1	T2	T3	SurrRev
1	A	2021	\$1,000	\$10,000
2	B	2021	\$0	\$5,000
3	C	2021	\$2,000	\$15,000

Technical Challenge: Double Counting

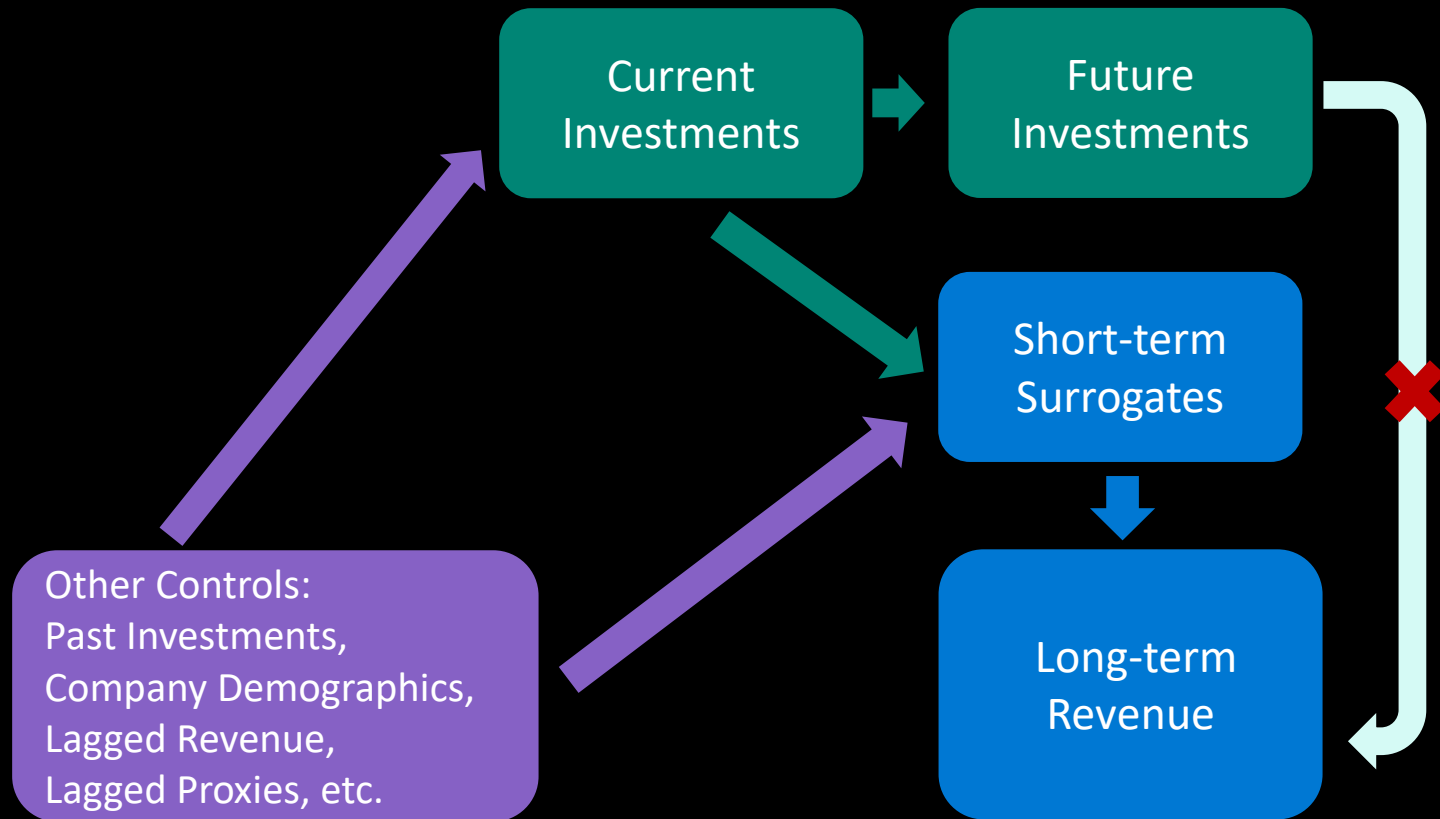
- Effect coming from future investments lead to **upwards biased effect**



- In the **historical dataset (0)**, long-term revenue includes the effect from future investment in the past, it will be **double attributed** to the current investments.

Methodology


- Double Machine Learning for dynamic effects (Lewis, Syrgkanis, "Double/Debiased Machine Learning for Dynamic Treatment Effects")



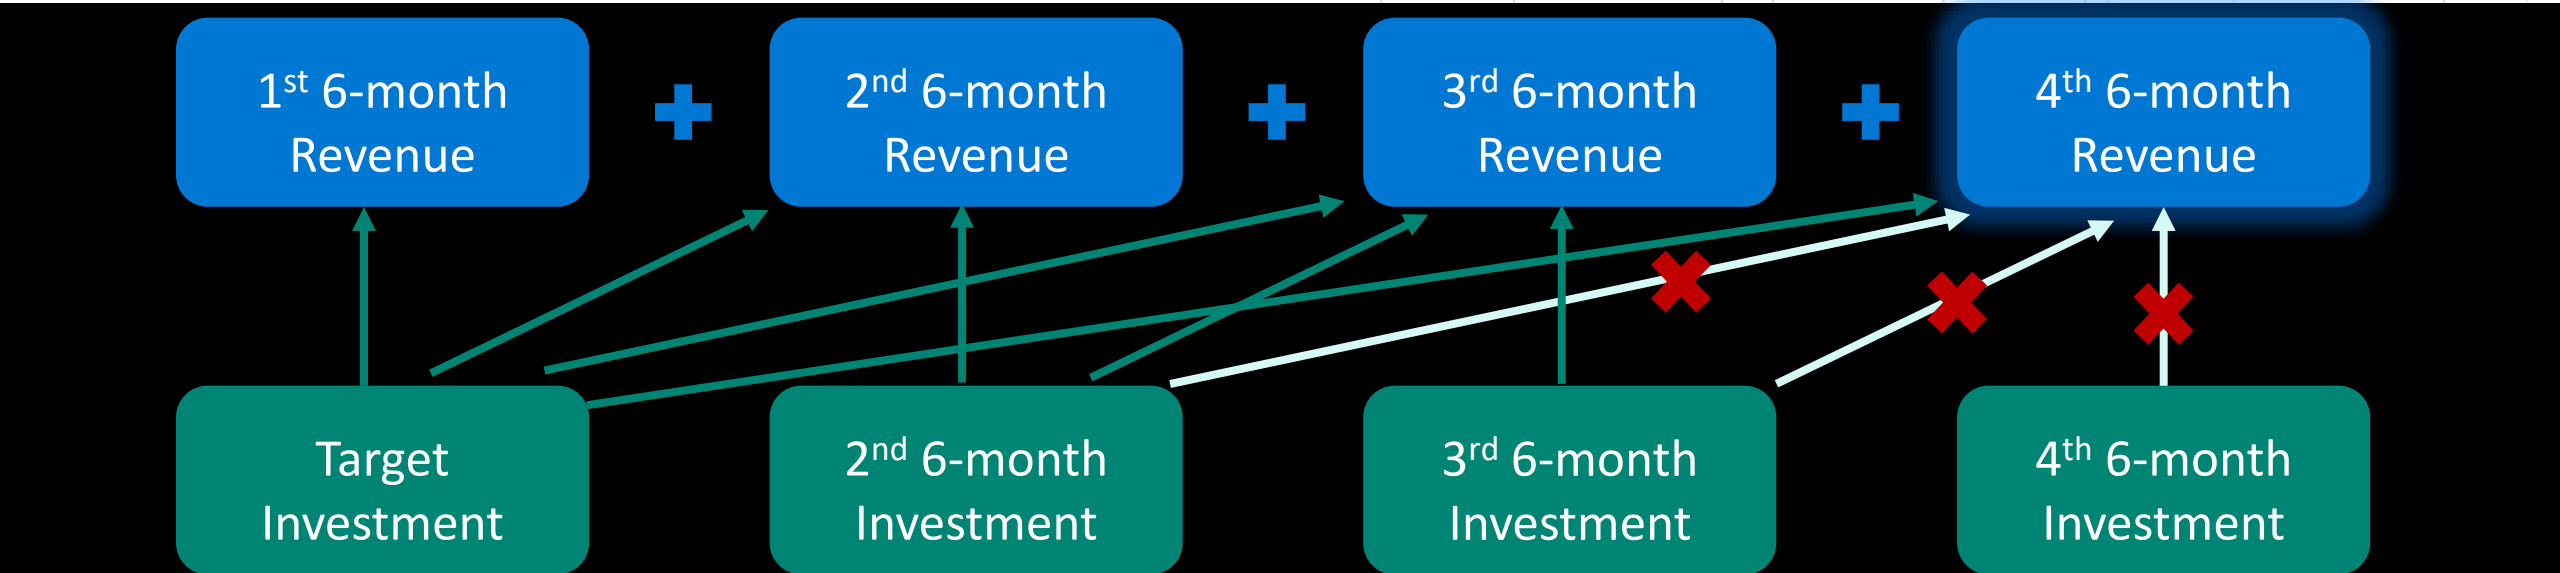
- **Dynamic treatment effect approach** removes the effects of future incentives from the **historical outcome** to create an adjusted long-term revenue as if those future incentives never happened.

Dynamic Double Machine Learning

- Start with simple 2 periods:
 - Run DML to estimate effect θ_{t+1} of T_{t+1} on Y_{t+1}
 - Subtract that effect to create the adjusted outcome
$$Y_{adj} = Y_t + (Y_{t+1} - \theta_{t+1} * T_{t+1})$$
- Use Y_{adj} as outcome to train surrogate model

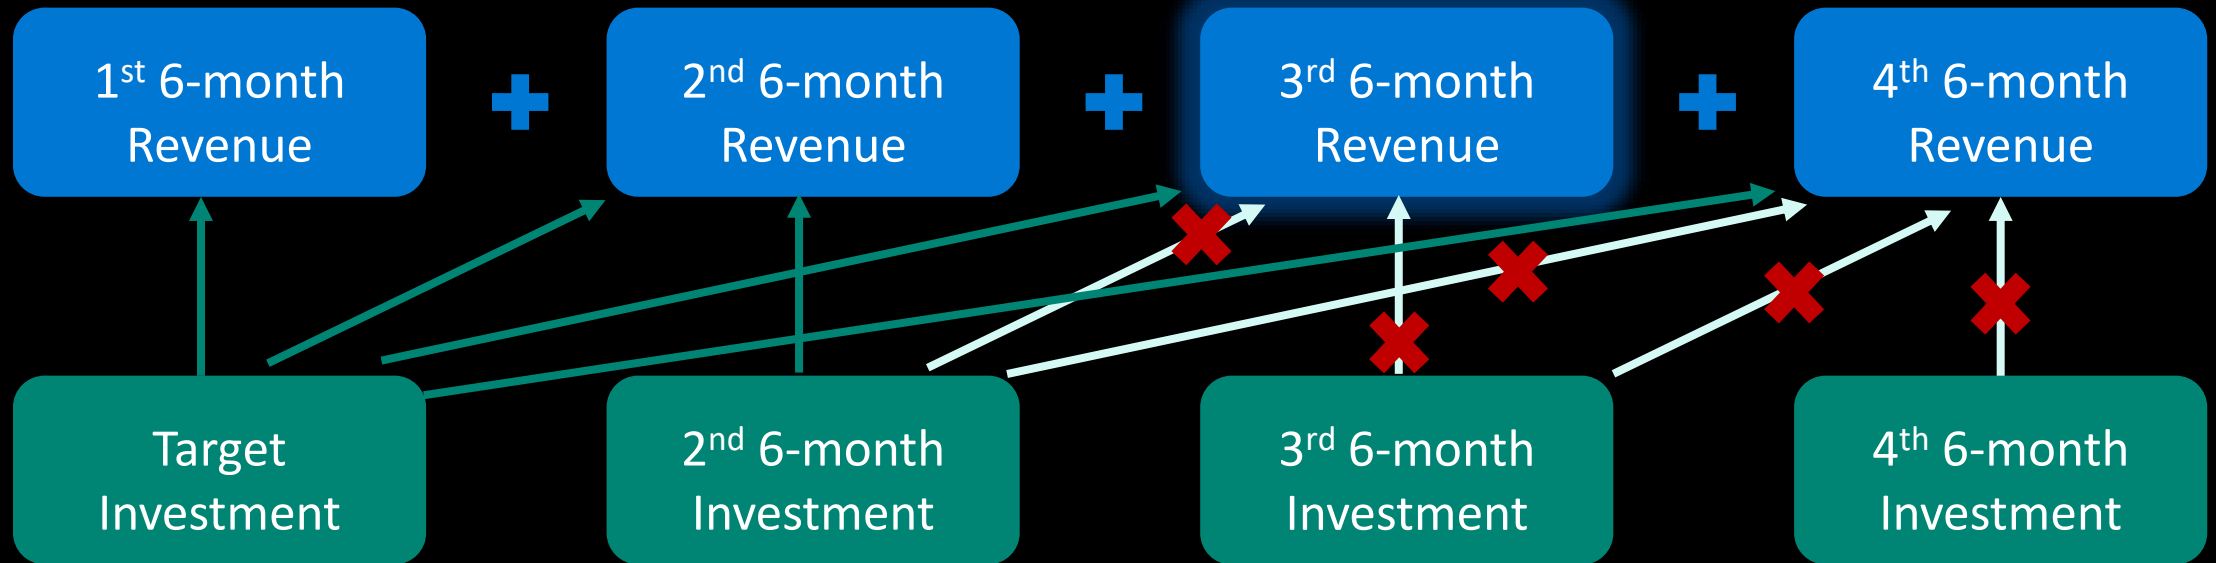


ent in ROI project



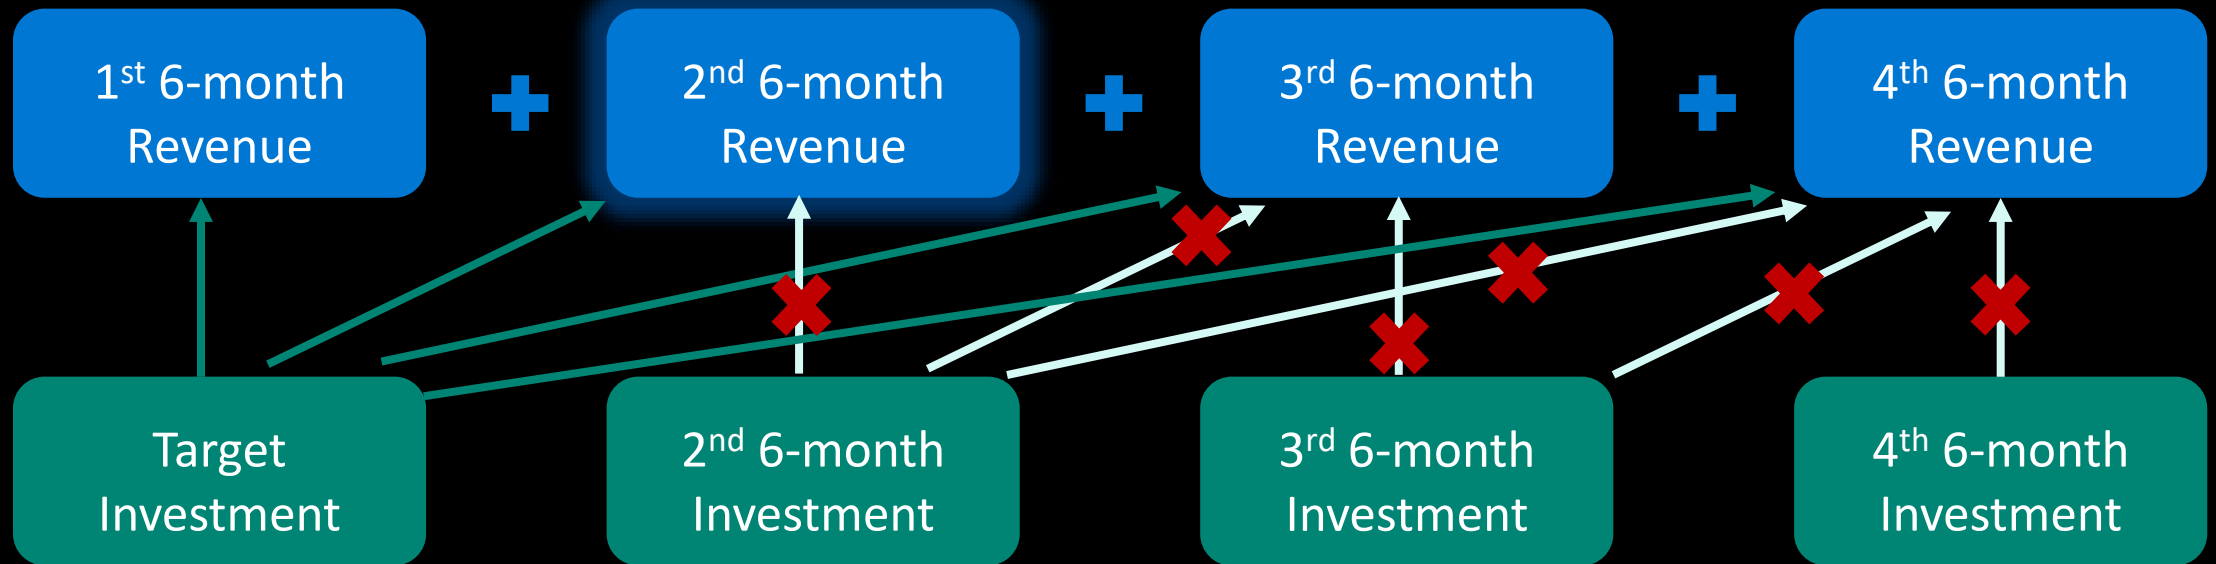
- From **historical dataset**, break the 2-year period into 6-month chunks
- **Recursively** subtract the effect of future investment from each revenue chunk
- **Sum up** the revenue chunks to get adjusted revenue, net of predicted effect of later investments
- Train surrogate index using **adjusted 2-year revenue**

Dynamic DML Adjustment in ROI project



- From **historical dataset**, break the 2-year period into 6-month chunks
- **Recursively** subtract the effect of future investment from each revenue chunk
- **Sum up** the revenue chunks to get adjusted revenue, net of predicted effect of later investments
- Train surrogate index using **adjusted 2-year revenue**

Dynamic DML Adjustment in ROI project



- From **historical dataset**, break the 2-year period into 6-month chunks
- **Recursively** subtract the effect of future investment from each revenue chunk
- **Sum up** the revenue chunks to get adjusted revenue, net of predicted effect of later investments
- Train surrogate index using **adjusted 2-year revenue**

EconML Solution for Dynamic DML

```
# print data shape
print("Outcome shape: ", panelY.shape)
print("Treatment shape: ", panelT.shape)
print("Controls shape: ", panelX.shape)
print("Group ID shape: ", panelGroups.shape)
```

```
Outcome shape: (5000, 4)
Treatment shape: (5000, 4, 3)
Controls shape: (5000, 4, 89)
Group ID shape: (5000, 4)
```

```
# imports
from econml.dml import DynamicDML

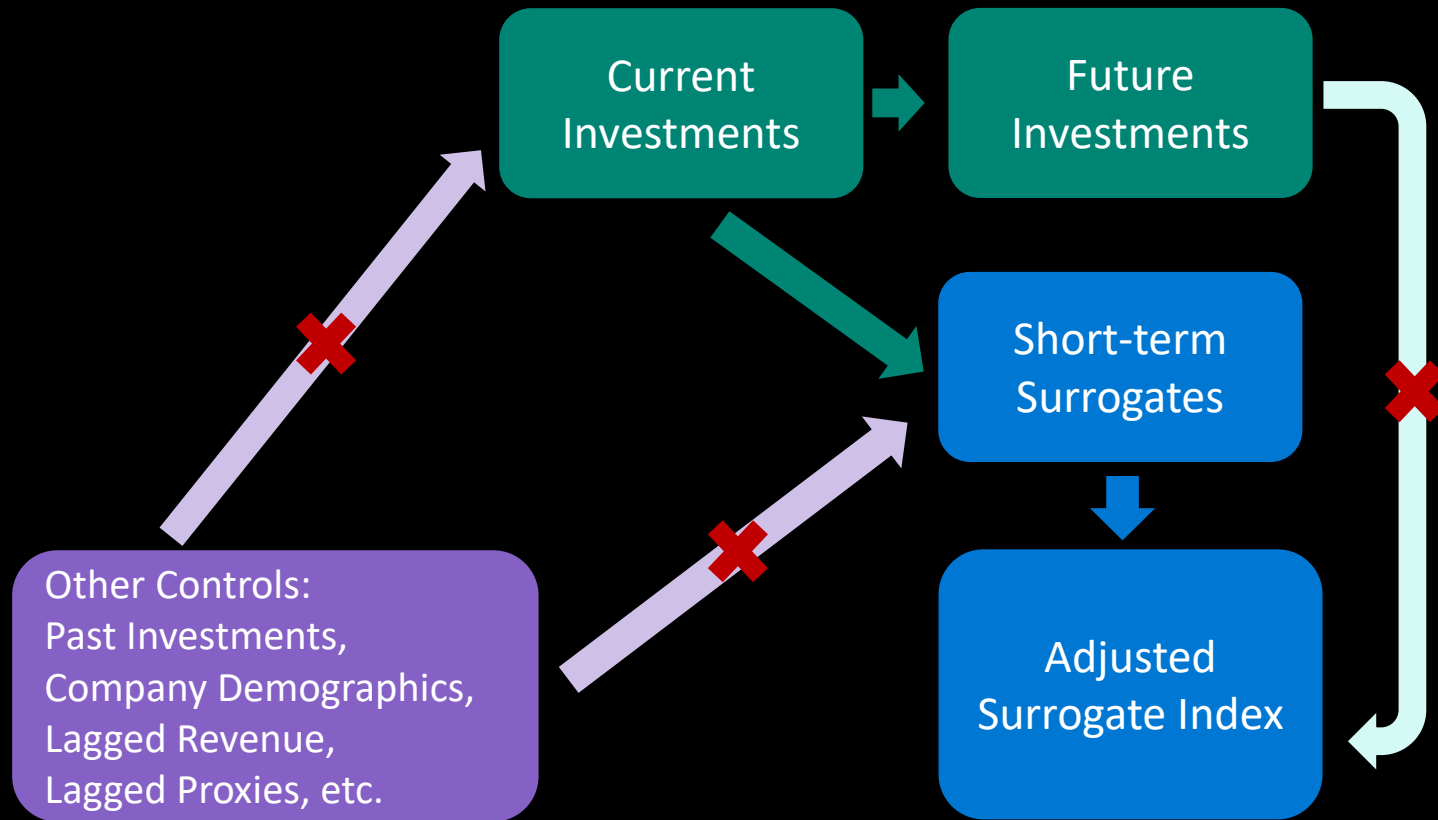
# initiate dynamic DML estimator
est = DynamicDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
# Learn period effect for each period T on last period revenue
est.fit(
    long(panelY), # reshape the panel data n = n_groups * n_periods
    long(panelT), # reshape the panel data n = n_groups * n_periods
    X=None,
    W=long(panelX), # reshape the panel data n = n_groups * n_periods
    groups=long(panelGroups), # reshape the panel data n = n_groups * n_periods
)
est.const_marginal_effect().reshape(-1, n_treatments)
```

```
array([[0.6897052 , 0.29836498, 0.1234395 ],
       [0.38710883, 0.1940497 , 0.09248946],
       [0.23237571, 0.0980073 , 0.03435173],
       [0.15665969, 0.05032955, 0.04140905]])
```

Return constant marginal effect for each period T on last period revenue for each treatment with shape (n_periods, d_t).

Updated Causal Graph

- Dynamic DML + Surrogate Index + DML

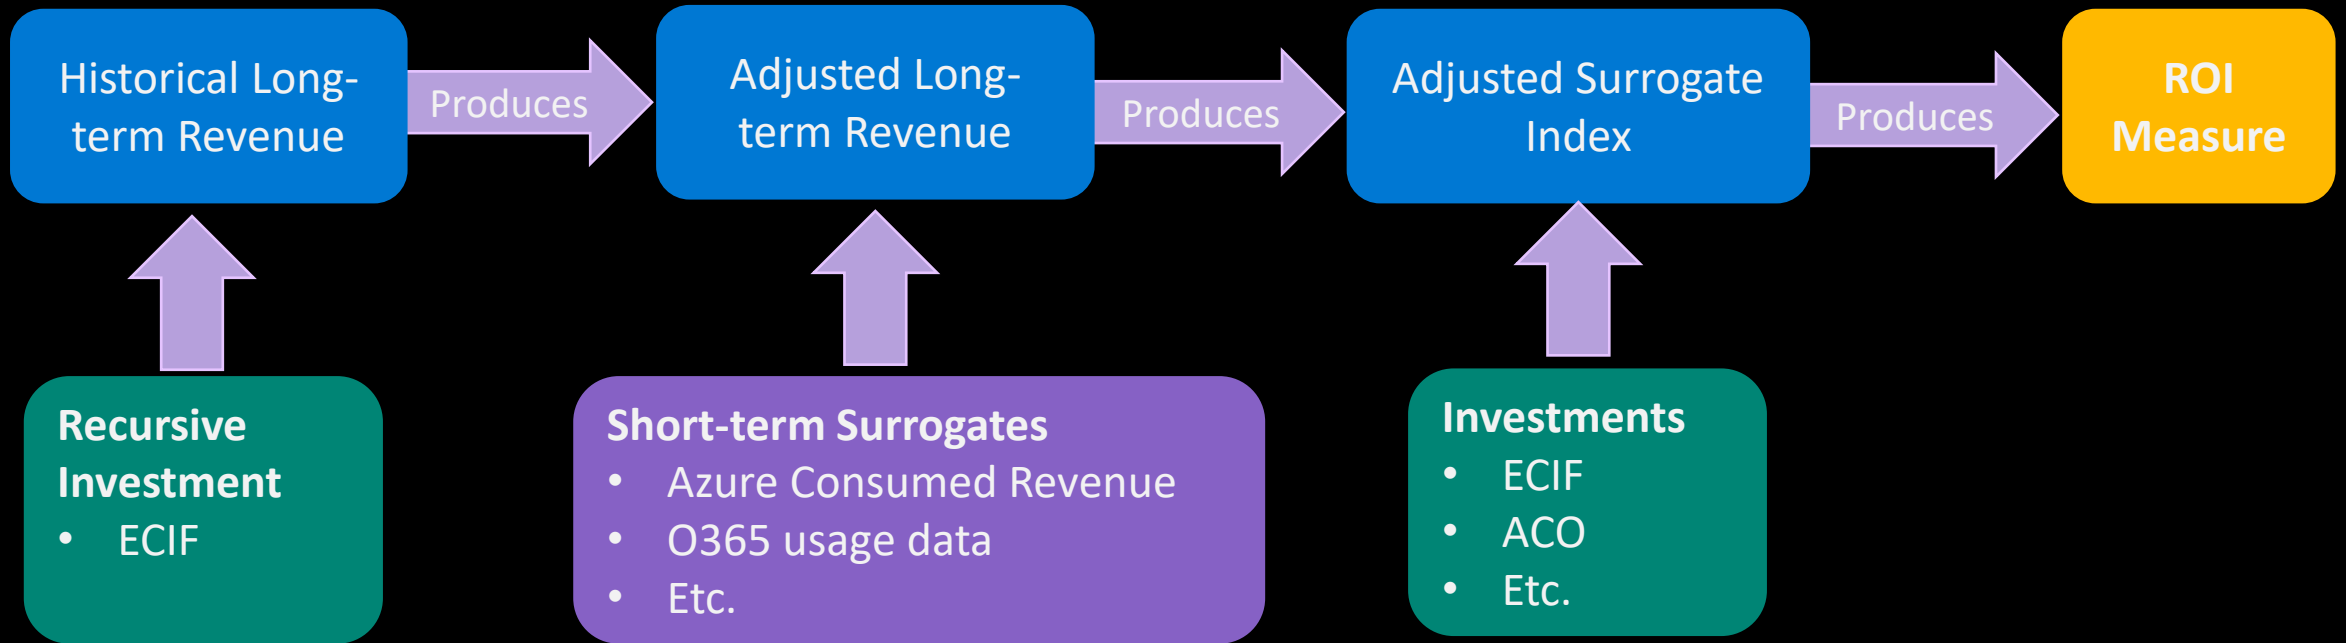


Unified Pipeline


Step 1: Adjust historical revenue by removing effect from future investment

Step 2: Forecast long-term revenue using short-term surrogates

Step 3: Estimate causal effect of investments on Adjusted Surrogate Index



EconML Solution



Historical Data

	Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2018	\$1,000	\$10,000
2	A	2019	\$2,000	\$12,000
3	A	2020	\$3,000	\$15,000
4	B	2018	\$0	\$5,000
5	B	2019	\$100	\$10,000
6	B	2020	\$1,200	\$7,000
7	C	2018	\$1,000	\$20,000
8	C	2019	\$1,500	\$25,000
9	C	2020	\$500	\$15,000

Current Data with New Investment

	Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2021	\$1,000	\$10,000
2	B	2021	\$0	\$5,000
3	C	2021	\$2,000	\$15,000

EconML Solution

Historical Data

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjRev
1	A	2018	...	\$1,000	\$10,000
2	A	2019	...	\$2,000	\$12,000
3	A	2020	...	\$3,000	\$15,000
4	B	2018	...	\$0	\$5,000
5	B	2019	...	\$100	\$10,000
6	B	2020	...	\$1,200	\$7,000
7	C	2018	...	\$1,000	\$20,000
8	C	2019	...	\$1,500	\$25,000
9	C	2020	...	\$500	\$15,000

```
# on historical data construct adjusted outcomes
from econml.dml import DynamicDML

est = DynamicDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
for t in range(1, n_periods): # for each target period 1...m
    est.fit(long(panelY_O[:, 1 : t + 1]), long(panelT_O[:, 1 : t + 1, :]), # reshape data to long format
            X=None, W=long(panelX_O[:, 1 : t + 1, :]), groups=long(panelGroups_O[:, 1 : t + 1]))
    # remove effect of observed treatments
    panelYadj_O[:, t] = panelY_O[:, t] - est.effect(T0=0, T1=wide(panelT_O[:, 1 : t + 1, :])) # reshape data to wide format
```

Current Data with New Investment

Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2021	...	\$1,000	\$10,000
2	B	2021	...	\$0	\$5,000
3	C	2021	...	\$2,000	\$15,000

EconML Solution

Historical Data

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjRev
1	A	2018	...	\$1,000	\$10,000
2	A	2019	...	\$2,000	\$12,000
3	A	2020	...	\$3,000	\$15,000
4	B	2018	...	\$0	\$5,000
5	B	2019	...	\$100	\$10,000
6	B	2020	...	\$1,200	\$7,000
7	C	2018	...	\$1,000	\$20,000
8	C	2019	...	\$1,500	\$25,000
9	C	2020	...	\$500	\$15,000

Current Data with New Investment

Company	Year	Features	Controls/Surrogates	T1	T2	T3	Revenue
1	A	2021	...	\$1,000	\$10,000
2	B	2021	...	\$0	\$5,000
3	C	2021	...	\$2,000	\$15,000

```
# on historical data construct adjusted outcomes
from econml.dml import DynamicDML

est = DynamicDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
for t in range(1, n_periods): # for each target period 1..m
    est.fit(long(panelY_0[:, 1 : t + 1]), long(panelT_0[:, 1 : t + 1, :]), # reshape data to long format
            X=None, W=long(panelX_0[:, 1 : t + 1, :]), groups=long(panelGroups_0[:, 1 : t + 1]))
    # remove effect of observed treatments
    panelYadj_0[:, t] = panelY_0[:, t] - est.effect(T0=0, T1=wide(panelT_0[:, 1 : t + 1, :])) # reshape data to wide format
```

```
# train surrogate index
XS_0 = np.hstack([panelX_0[:, 0], panelYadj_0[:, :1]]) # concatenate controls and surrogates from historical dataset
TotalYadj_0 = np.sum(panelYadj_0, axis=1) # total revenue from historical dataset
adjusted_proxy_model = LassoCV().fit(XS_0, TotalYadj_0) # train proxy model from historical dataset
```

EconML Solution

Historical Data

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjRev
1	A	2018	...	\$1,000	\$10,000
2	A	2019	...	\$2,000	\$12,000
3	A	2020	...	\$3,000	\$15,000
4	B	2018	...	\$0	\$5,000
5	B	2019	...	\$100	\$10,000
6	B	2020	...	\$1,200	\$7,000
7	C	2018	...	\$1,000	\$20,000
8	C	2019	...	\$1,500	\$25,000
9	C	2020	...	\$500	\$15,000

Current Data with New Investment

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjSurrRev
1	A	2021	...	\$1,000	\$10,000
2	B	2021	...	\$0	\$5,000
3	C	2021	...	\$2,000	\$15,000

```
# on historical data construct adjusted outcomes
from econml.dml import DynamicDML

est = DynamicDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
for t in range(1, n_periods): # for each target period 1..m
    est.fit(long(panelY_0[:, 1 : t + 1]), long(panelT_0[:, 1 : t + 1, :]), # reshape data to long format
            X=None, W=long(panelX_0[:, 1 : t + 1, :]), groups=long(panelGroups_0[:, 1 : t + 1]))
    # remove effect of observed treatments
    panelYadj_0[:, t] = panelY_0[:, t] - est.effect(T0=0, T1=wide(panelT_0[:, 1 : t + 1, :])) # reshape data to wide format
```

```
# train surrogate index
XS_0 = np.hstack([panelX_0[:, 0], panelYadj_0[:, :1]]) # concatenate controls and surrogates from historical dataset
TotalYadj_0 = np.sum(panelYadj_0, axis=1) # total revenue from historical dataset
adjusted_proxy_model = LassoCV().fit(XS_0, TotalYadj_0) # train proxy model from historical dataset
```

```
# predict current long term revenue
XS_E = np.hstack([panelX_E[:, 0], panelY_E[:, :1]]) # concatenate controls and surrogates from current dataset
sindex_adj = adjusted_proxy_model.predict(XS_E)
```

EconML Solution

Historical Data

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjRev
1	A	2018	...	\$1,000	\$10,000
2	A	2019	...	\$2,000	\$12,000
3	A	2020	...	\$3,000	\$15,000
4	B	2018	...	\$0	\$5,000
5	B	2019	...	\$100	\$10,000
6	B	2020	...	\$1,200	\$7,000
7	C	2018	...	\$1,000	\$20,000
8	C	2020	...	\$1,500	\$25,000
9	C	2020	...	\$500	\$15,000

```
# on historical data construct adjusted outcomes
from econml.dml import DynamicDML

est = DynamicDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
for t in range(1, n_periods): # for each target period 1..m
    est.fit(long(panelY_O[:, 1 : t + 1]), long(panelT_O[:, 1 : t + 1, :]), # reshape data to long format
            X=None, W=long(panelX_O[:, 1 : t + 1, :]), groups=long(panelGroups_O[:, 1 : t + 1]))
    # remove effect of observed treatments
    panelYadj_O[:, t] = panelY_O[:, t] - est.effect(T0=0, T1=wide(panelT_O[:, 1 : t + 1, :])) # reshape data to wide format
```

```
# train surrogate index
XS_O = np.hstack([panelX_O[:, 0], panelYadj_O[:, :1]]) # concatenate controls and surrogates from historical dataset
TotalYadj_O = np.sum(panelYadj_O, axis=1) # total revenue from historical dataset
adjusted_proxy_model = LassoCV().fit(XS_O, TotalYadj_O) # train proxy model from historical dataset
```

```
# predict current long term revenue
XS_E = np.hstack([panelX_E[:, 0], panelY_E[:, :1]]) # concatenate controls and surrogates from current dataset
sindex_adj = adjusted_proxy_model.predict(XS_E)
```

```
# Learn treatment effect on surrogate index
from econml.dml import LinearDML
est = LinearDML(model_y=LassoCV(max_iter=1000), model_t=MultiTaskLassoCV(max_iter=1000))
# fit treatment_0 on total revenue from current dataset
est.fit(sindex_adj, panelT_E[:, 0], X=None, W=panelX_E[:, 0])
# print treatment effect summary
est.summary(alpha=0.05)
```

Current Data with New Investment

Company	Year	Features	Controls/Surrogates	T1	T2	T3	AdjSurrRev
1	A	2021	...	\$1,000	\$10,000
2	B	2021	...	\$0	\$5,000
3	C	2021	...	\$2,000	\$15,000

CATE Intercept Results

	point_estimate	stderr	zstat	pvalue	ci_lower	ci_upper
cate_intercept T0	1.284	0.024	52.626	0.0	1.236	1.332
cate_intercept T1	0.603	0.011	53.221	0.0	0.581	0.626
cate_intercept T2	0.242	0.005	46.863	0.0	0.232	0.253

Wrap Up

- **Surrogate index approach** is a widely useful technique for measuring **long-term** effects from **short-term** data (healthcare, business, marketing)
- Typical assumptions for the method to work can **be severely violated** if historical data contain other treatments
- Adaptivity and auto-correlation of the historical treatment policy can severely **bias effects**
- These biases can be corrected by combining techniques from **estimation of treatment effects in the dynamic treatment regime with the surrogate approach**
- If careful in the estimation method (Neyman orthogonal moments + sample splitting), **Machine learning** can be used to enable estimation with **high-dimensional** surrogates or controls

Applying EconML to Your Causal Problem

- Jupyter notebooks for similar use cases:
 - <https://github.com/microsoft/EconML/tree/master/notebooks/CustomScenarios>
- Learn more about all the uses of EconML on our website
 - <https://aka.ms/econml>
- Github and Doc link:
 - <https://github.com/Microsoft/EconML>
 - <https://econml.azurewebsites.net/>
- Get started with EconML in Python: `pip install econml`

Reference

- [1] Susan Athey, Raj Chetty, Guido Imbens, and Hyunseung Kang. Estimating treatment effects using multiple surrogates: The role of the surrogate score and the surrogate index, 2020.
- [2] Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins. Double/debiased machine learning for treatment and structural parameters. The Econometrics Journal, 21(1):C1–C68, 01 2018.
- [3] Greg Lewis and Vasilis Syrgkanis. Double/debiased machine learning for dynamic treatment effects, 2020.
- [4] Keith Battocchi, Eleanor Dillon, Maggie Hei, Greg Lewis, Miruna Oprescu, Vasilis Syrgkanis. Estimating the long-term effects of novel treatments, 2021.
- [5] Daniel Yehdego, Jane Huang, Saurabh Kumar, and Siddharth Kumar. An application of causal modeling for azure investment attribution, 2020.

QA

